

```
=====
TXT2PDF V1.7.1 - THE SPARTAN PDF GENERATOR
Zero-dependency TXT to PDF converter in pure C11
=====
```

"Turn plain text into structured PDFs without asking for permission."

```
Language:  C11 (pure, no C++ runtime, no external libs)
Binary:    Static, stripped, musl libc, tiny
Output:    PDF 1.4 with outlines, internal links, Base14 fonts, images
Status:    v1.7.1 feature release
```

```
=====
QUICK START
=====
```

- Compile with the provided script:

```
./compile.sh
```

- Compile manually with musl:

```
musl-gcc -Os -Wall -Wextra -std=c11 -static -o txt2pdf txt2pdf.c
strip txt2pdf
```

- Run with default Courier font:

```
./txt2pdf readme.txt readme.pdf
```

- Run with Helvetica:

```
./txt2pdf --sans readme.txt readme.pdf
```

- Run with Times-Roman:

```
./txt2pdf --serif readme.txt readme.pdf
```

- Run with debug diagnostics:

```
./txt2pdf --debug readme.txt readme.pdf
```

```
=====
[TOC] TABLE OF CONTENTS
=====
```

1. what is txt2pdf?
2. what is new in v1.7.1?
3. fonts and typography
4. image support
5. parser rules
6. table of contents rules
7. usage and debug mode
8. minimal working structure
9. pdf internals
10. project structure
11. building and static cookbook
12. known limitations
13. troubleshooting
14. design philosophy
15. license and credits

1. WHAT IS TXT2PDF?

txt2pdf is a zero-dependency command-line utility that converts plain text files into structured PDF 1.4 documents.

Unlike Pandoc, wkhtmltopdf, or Python-based converters that require hundreds of megabytes of runtime dependencies, txt2pdf is a small static binary.

It parses text, calculates page breaks, generates PDF objects, and writes the final binary stream directly to disk using only the C standard library.

It is designed for environments where installing dependencies is impossible, unwanted, or where the network is down.

2. WHAT IS NEW IN V1.7.1?

Version 1.7.1 adds better winansi char detection

Version 1.7 adds the two most requested Spartan features:

- Base14 font selection.
- Standalone block images.

Main changes:

- Added --sans, --serif and --mono options.
- Added Helvetica and Times-Roman Base14 font support.
- Added proportional text wrapping using built-in ASCII metrics.
- Added [IMG]...[/IMG] standalone image support.
- Added JPEG passthrough using native PDF DCTDecode.
- Added simple PNG passthrough using native PDF FlateDecode.
- Added debug diagnostics for detected images.
- Kept the project dependency-free and single-file friendly.

The default behavior remains unchanged:

- Courier is still the default font.
- Existing v1.6 documents should convert exactly as before.
- Images and proportional fonts are opt-in features.

3. FONTS AND TYPOGRAPHY

txt2pdf v1.7 supports the standard PDF Base14 font families.

Available options:

- --mono
Use Courier. This is the default and the classic txt2pdf layout.
- --sans
Use Helvetica. This gives a clean, modern, proportional look.
- --serif
Use Times-Roman. This gives a classic documentation look.

Examples:

```
./txt2pdf --mono input.txt output.pdf
./txt2pdf --sans input.txt output.pdf
./txt2pdf --serif input.txt output.pdf
```

Important notes:

- Base14 fonts are not embedded into the PDF.
- PDF viewers normally provide built-in substitutes for Base14 fonts.
- Courier layout remains fixed-width and highly predictable.
- Helvetica and Times-Roman use built-in ASCII metrics for line wrapping.
- Accents and non-ASCII characters remain limited to WinAnsiEncoding.

4. IMAGE SUPPORT

txt2pdf v1.7 supports standalone block images.

Supported image tags:

- [IMG]photo.jpg[/IMG]
- [IMG]diagram.jpeg[/IMG]
- [IMG]chart.png[/IMG]

The image line must be the only content on that line.

Good image line:

- [IMG]chart.png[/IMG]

Bad image line:

- See this chart: [IMG]chart.png[/IMG]

Supported formats:

- JPEG
Embedded using PDF native DCTDecode passthrough.
- PNG
Embedded using a minimalist passthrough for simple PNG files.

Supported PNG types:

- 8-bit grayscale
- 8-bit RGB
- 8-bit palette

Unsupported PNG types:

- alpha transparency
- interlaced PNGs
- 16-bit channels
- palette transparency

Image layout behavior:

- Images are treated as standalone content blocks.
- Aspect ratio is preserved.
- Images wider than the text area are scaled down.

- Images taller than the page are scaled down to fit.
- An image may force a page break if it does not fit.
- Image paths are relative to the current working directory.
- The PDF stores the original compressed image data.
- Reducing source image size reduces final PDF size.

5. PARSER RULES

The parser creates PDF bookmarks only for lines that are real section headers.

A line is considered a section header only if all of these conditions are true:

- It starts with a digit, optionally followed by letters, then a dot.
- The title text after the dot contains only uppercase letters.
- The line is surrounded by separator lines made mostly of '=' characters.
- The line is not inside a contents block.

Valid section shape:

- separator line
- uppercase numbered title
- separator line

Rejected lines:

- numbered titles containing lowercase letters
- numbered titles used as examples inside normal paragraphs
- numbered titles not surrounded by separator lines
- numbered titles inside the contents block

This makes the parser much safer for manuals that contain examples.

The parser supports hierarchical IDs such as:

- 1
- 2.3
- 2.3.1
- 11.14
- 2.0a
- 13.1b
- 26.5a

6. TABLE OF CONTENTS RULES

Internal links are created only from contents blocks.

A contents block is detected when:

- a line contains the token "[TOC]"
- the next non-blank line is a separator line
- the block ends at the next separator line

Rules for entries:

- entries must be numbered lines

- lowercase entries are recommended
- entries are not converted into bookmarks
- entries are linked to the matching section by number or title
- uppercase entries inside the contents block are also ignored as bookmarks

This means the old duplicate-bookmark problem is largely eliminated by the parser itself.

=====

7. USAGE AND DEBUG MODE

=====

Basic usage:

```
./txt2pdf input.txt output.pdf
```

Font selection:

```
./txt2pdf --mono input.txt output.pdf
./txt2pdf --sans input.txt output.pdf
./txt2pdf --serif input.txt output.pdf
```

Debug usage:

```
./txt2pdf --debug input.txt output.pdf
```

Combined usage:

```
./txt2pdf --debug --sans input.txt output.pdf
```

Debug mode prints diagnostics to stderr.

It reports:

- detected encoding
- selected font
- lines marked as part of a contents block
- detected section headers
- skipped uppercase numbered lines
- detected internal links
- detected images
- image placements

Use debug mode when:

- a bookmark is missing
- too many bookmarks appear
- a contents entry does not link anywhere
- an image does not appear
- you are adapting a new readme to the Spartan format

=====

8. MINIMAL WORKING STRUCTURE

=====

A minimal compatible document needs three things:

- a contents block starting with the [TOC] token
- lowercase or mixed-case contents entries
- uppercase section headers surrounded by separator lines

Recommended structure:

- write the contents title line containing [TOC]
- write a separator line
- write entries such as "1. introduction" and "2. conclusion"
- write a separator line
- for each section, write a separator line
- write the uppercase header, for example "1. INTRODUCTION"
- write another separator line
- write the section body

Important:

Do not wrap example headers with separator lines unless you want them to become real bookmarks.

```
=====
9. PDF INTERNALS
=====
```

txt2pdf generates a valid PDF 1.4 structure from scratch.

Catalog and Pages:

Calculates A4 dimensions, applies margins, wraps text, and places content using a line-based layout.

Fonts:

Generates a Base14 Type1 font object. The default is Courier. The --sans and --serif options switch to Helvetica or Times-Roman.

Images:

Loads image metadata, reserves vertical layout space, writes image XObject streams, and references them from page resources.

Encoding and Smart Decoder:

Auto-detects UTF-8 vs Latin-1. If UTF-8 is detected, a smart decoder translates supported multi-byte sequences to WinAnsiEncoding.

Outlines:

Generates a linked list of PDF Outline objects pointing to the exact XYZ coordinates of each uppercase header.

Annotations:

Scans contents blocks and injects link annotations for each entry matching a real section.

Cross-Reference Table:

Calculates byte offsets for every object and writes a compliant xref table and trailer.

```
=====
10. PROJECT STRUCTURE
=====
```

The txt2pdf codebase follows a single-file architecture.

Typical layout:

- txt2pdf.c
- compile.sh
- readme.txt

The main C file contains:

- argument parsing
- encoding detection
- text parsing
- header detection
- contents detection
- image loading
- font selection
- page break calculation
- PDF object generation
- binary stream writing

The tool is intentionally small and auditable.

```
=====
11. BUILDING AND STATIC COOKBOOK
=====
```

Provided build script:

```
./compile.sh
```

Dynamic build:

```
gcc -O2 -Wall -Wextra -std=c11 -o txt2pdf txt2pdf.c
```

Static musl build:

```
musl-gcc -Os -Wall -Wextra -std=c11 -static -o txt2pdf txt2pdf.c
strip txt2pdf
```

Why musl?

glibc static binaries are often much larger due to NSS and locale data.
musl provides a clean, minimal libc that results in a tiny, fully
self-contained executable.

Verify static build:

```
file txt2pdf
ldd txt2pdf
```

Expected:

- file reports statically linked
- ldd says it is not a dynamic executable

```
=====
12. KNOWN LIMITATIONS
=====
```

- No external URL hyperlinks.

- No custom page sizes.
- No custom margins.
- No multi-column layout.
- No inline images.
- No image captions.
- No repeated header/footer images.
- No per-image width, height, DPI, or centering attributes yet.
- PNG support is limited to simple 8-bit non-interlaced images.
- PNG alpha transparency is not supported.
- Interlaced PNGs are not supported.
- WinAnsiEncoding only.
- Characters outside WinAnsi are replaced with '?'.
 - Section headers must be fenced by separator lines.
 - Internal links are only generated from contents blocks.
 - Base14 fonts are not embedded.

13. TROUBLESHOOTING

Missing bookmarks:

- The header contains lowercase letters.
- The header is not surrounded by separator lines.
- The header is accidentally inside a contents block.

Duplicate bookmarks:

- Old documents may have uppercase contents entries.
- The updated parser ignores contents entries as sections.
- If duplicates persist, run with `--debug` and inspect skipped lines.

Missing links:

- The entry is not inside a detected contents block.
- The entry number does not match any section.
- The entry title does not match any section title.

Extra bookmarks:

- An example header is surrounded by separator lines.
- Remove the separators around examples or rewrite the example line.

Garbled characters:

- Ensure your text editor saves the file in UTF-8.
- Run with `--debug` to verify the detected encoding.
- Characters outside WinAnsiEncoding cannot be rendered.

Image does not appear:

- The file path is wrong.
- The image line contains extra text beside the tag.
- The image format is not supported.
- The PNG uses alpha, interlacing, or 16-bit channels.

Image is too large:

- v1.7 does not yet support per-image width or DPI attributes.
- Use a smaller source image or crop it.
- High-resolution images are scaled visually but still embedded.

Font does not look as expected:

- Use --mono for the classic Courier layout.
- Use --sans for Helvetica.
- Use --serif for Times-Roman.
- Some viewers substitute Base14 fonts with similar fonts.

Debug output is the fastest way to diagnose the problem.

=====

14. DESIGN PHILOSOPHY

=====

Modern software development has forgotten how to build small tools.

We wrap simple text transformations in Electron apps, Node.js runtimes, and Docker containers.

txt2pdf is a rejection of that bloat.

It proves that a fully functional, structured document generator can fit in a tiny static binary, written in pure C11, with zero external dependencies.

When the cloud goes down, and the package managers fail to resolve, the Spartan tool remains.

"Code is neither created nor destroyed; it is distilled."

=====

15. LICENSE AND CREDITS

=====

Source code : GPLv3

Manual: GFDL v1.3

txt2pdf was written as a demonstration of minimalist software design.

Inspired by the Unix philosophy: do one thing well.

Built with minimalism, on a potato, for potatoes.

=====

END OF DOCUMENTATION

=====